

# Rejection Sampling, Accept-Reject Method

---

Mark M. Fredrickson ([mfredric@umich.edu](mailto:mfredric@umich.edu))

Computational Methods in Statistics and Data Science (Stats 406)

# Rejection Sampling

---

## Conditioning

When generating **Laplace** RVs we used **conditioning** by noting if

$$f(x) = \frac{1}{2}e^{-|x|}$$

then

$$X \mid X > 0 \sim \text{Exp}(1).$$

The reverse could also be useful, if we had a **a source of Laplace RVs** we could use it to **generate Exp(1)**.

Let's illustrate this by first getting the **quantile function** and then using the **inversion method** to make `rlaplace`.

$f(x)$  to  $F(x)$

As we did before, let's split up  $f(x) = (1/2)e^{-|x|}$  when  $x < 0$  and  $x \geq 0$ :

$$f(x) = \begin{cases} \frac{1}{2}e^x & : x < 0 \\ \frac{1}{2}e^{-x} & x \geq 0 \end{cases}$$

Then we have a **piece-wise CDF**. For  $x < 0$ ,

$$F(x) = \frac{1}{2} \left( \int_{-\infty}^x e^t dt \right) = \frac{1}{2}e^x$$

For  $x \geq 0$  we can decompose  $F(x)$  as:

$$F(x) = F(0) + P(0 \leq X \leq x) = \frac{1}{2} + \frac{1}{2} \left( \int_0^x e^{-t} dt \right) = 1 - \frac{1}{2}e^{-x}$$

The quantile function will also be **piece-wise** with the change at:

$$F(0) = 1/2 \Rightarrow Q(1/2) = 0$$


Solving the CDF leads to

$$F(x) = \begin{cases} \frac{1}{2}e^x & : x < 0 \\ 1 - \frac{1}{2}e^{-x} & : x \geq 0 \end{cases} \Rightarrow Q(u) = \begin{cases} \log(2u) & : 0 \leq u < \frac{1}{2} \\ \log\left(\frac{1}{2-2u}\right) & : \frac{1}{2} \leq u \leq 1 \end{cases}$$

$$Q(u) = \begin{cases} \log(2u) : 0 \leq u < \frac{1}{2} \\ \log\left(\frac{1}{2-2u}\right) : \frac{1}{2} \leq u \leq 1 \end{cases}$$

```
> rlaplace <- function(n) {  
+   u <- runif(n)  
+   ifelse(u < 1/2, log(2 * u), log(1 / (2 - 2 * u)))  
+ }
```

## Density plot



## Rejection Sampling for $\text{Exp}(1)$

Suppose we have `rlaplace` but not `rexp`. We saw that if  $X \sim \text{Laplace}(0)$ , then  $X \mid X > 0 \sim \text{Exp}(1)$ . Let's do that:

```
> x <- rlaplace(1000)
> x_positive <- keep(x, x > 0)
> mean(x_positive) # should be close to 1

[1] 0.9407

> t.test(x_positive, conf.level = 0.999)$conf.int

[1] 0.8021 1.0793
attr(,"conf.level")
[1] 0.999
```



## Variance of the Estimator

The method worked fairly well, but **had to throw away about 50% of RVs:**

```
> length(x_positive)
```

```
[1] 485
```

Recall the variance of a the sample mean is:

$$\text{Var}(\bar{X}) = \frac{1}{n} \text{Var}(X)$$

by throwing away samples, we have **increased the variance** (relative to having a way to keep all the samples).

## Rejection Sampling in General

The idea of **rejection sampling** can be used when our **target RV** can be expressed as a subset of another **candidate RV**.

Examples:

- Conditioning on the variable itself,  $X \sim \text{Poisson}(\lambda)$ ,  $Y = X \mid X \text{ is odd}$
- Truncated distributions, e.g.  $X \sim N(0, 1)$ ,  $Y = X \mid a < X < b$
- Uniform points on the unit circle:  $U_1, U_2 \sim U(-1, 1)$  (independent),  
 $(V_1, V_2) = (U_1, U_2) \mid U_1^2 + U_2^2 \leq 1$

Procedure: generate from candidate, only keep results that meet criteria.

## Example: Truncated Normal Distribution

A **truncated distribution** takes a given distribution and **limits the support**.


For example,  $X \sim N(0, 1)$  and  $Y = X \mid 0 \leq X \leq 1$ . We can sample by drawing from  $X$  and only keeping those that fall in  $(0, 1)$ .

```
> x <- rnorm(10000)
> y <- keep(x, 0 < x & x < 1)
> length(y) # number of samples kept

[1] 3390

> mean(y) # use the samples to estimate E(Y)

[1] 0.4621
```



Consider the following method for approximating the value of  $\pi$ :


- Initialize  $a = 0$
- Generate  $U_1$  and  $U_2$  (independent) from  $U(0, 1)$
- If  $\sqrt{U_1^2 + U_2^2} \leq 1$  then  $a = a + 1$
- Repeat  $k$  times.
- Approximate  $\pi$  as

$$\pi = \frac{4a}{k}$$

## Calculating in R

```
> approx_pi <- function(k) {  
+   u1 <- runif(k)  
+   u2 <- runif(k)  
+   sqs <- sqrt(u1^2 + u2^2)  
+   return(4 * sum(sqs <= 1) / k)  
+ }
```

## Approximation vs. $k$



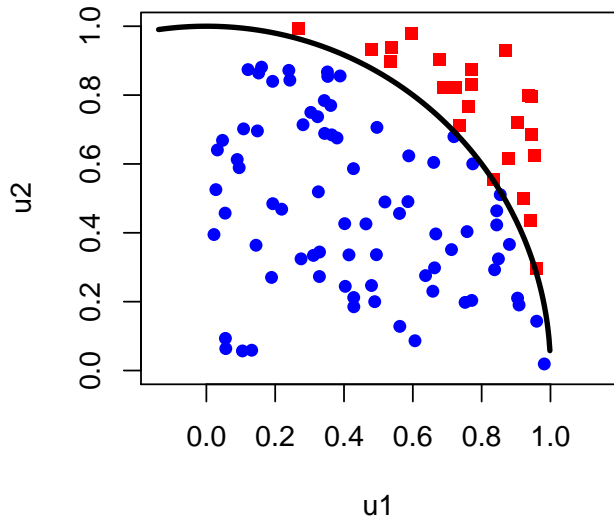
## Why does it work?

Recall: the circle centered at  $(0,0)$  and with radius 1, has area  $\pi$ .

We “throw darts” at the upper quadrant of the circle and see how many hit the circle.

This proportion should be about  $\pi/4$





## **Accept-Reject Method**


---


## Another interpretation of Approximating $\pi$


Notice that the following is a valid **probability density function**:

$$f(x) = \frac{4}{\pi} \sqrt{1 - x^2}, 0 \leq x \leq 1$$

This curve is just a **scaled version of the unit circle**.







## Interpretation

- We started by drawing  $U_1$  and  $U_2$ , and then keeping  $U_2$  if  $U_1^2 + U_2^2 \leq 1$ .
- This is equivalent to picking  $U_2$ , and keeping  $U_1$  if

$$\frac{4}{\pi} U_2 \leq \frac{4}{\pi} \sqrt{1 - U_1^2} = f(U_1)$$

or written another way:

$$U_2 \leq \frac{\pi f(U_1)}{4 g(U_1)}, \quad g(x) = 1 \text{ (pdf of } U_1)$$

- This process generated samples from  $X$  (the variable with density  $f(x)$ )!

## Accept-Reject in General

Suppose we want to sample from density (or mass) function:

$$f(x)$$

We know  $f(x)$ , but can't easily sample from it directly (e.g., Normal distribution).

But what if we had **another density**  $g(y)$  such that

$$c \times \frac{g(x)}{f(x)} \geq 1$$

for some  $c > 0$  and for all  $x$  such that  $f(x) > 0$ .

In other words we need  $cg(x)$  to lie above  $f(x)$  for any  $x$  where  $f(x)$  is positive.



## Accept-reject for $f(x)$

For the example we just did the **target density** was

$$f(x) = \frac{4}{\pi} \sqrt{1 - x^2}, 0 \leq x \leq 1$$

The **candidate density** was the uniform distribution  $g(x) = 1, 0 \leq x \leq 1$

The **constant** was

$$c = \frac{4}{\pi}$$

This ensures that:

$$cg(x) = \frac{4}{\pi} \geq \frac{4}{\pi} \sqrt{1 - x^2} = f(x)$$

(NB:  $c$  is not unique. We could have pick any  $c \geq 4/\pi$ .)

## Example: Uniform and Beta

Suppose we want to sample from the Beta(2,2) distribution which is:

$$f(x) = 6x(1 - x), \quad 0 \leq x \leq 1$$


but we only have the standard uniform  $g(x) = 1$ .

Observe that  $f(x)$  achieves its max at  $x = 0.5$ , the maximum value is  $f(0.5) = 6/4$ .

Therefore we have

$$\frac{6}{4} \frac{g(x)}{f(x)} \geq 1$$

for any point in  $[0, 1]$ .



# Accept-Reject Algorithm

- Draw a  $Y$  from  $g(y)$
- Draw  $U \sim U(0, 1)$
- If

$$U < \frac{f(Y)}{cg(Y)}$$

then **accept**  $X = Y$


- Otherwise, **reject**  $Y$  as a candidate and repeat.

## Example: Beta/Uniform

Say we generate  $Y = 0.3$ .

- $f(0.3) = 1.26$
- $g(0.3) = 1$
- We will accept if

$$U < \frac{1.26}{(6/4) \times 1} = 0.84$$




## R implementation


```
> k <- 10000
> ys <- runif(k)
> fys <- 6 * ys * (1 - ys)
> gys <- 1
> ratios <- fys / (gys * (6/4))

> us <- runif(k)
> accept <- us < ratios
> accepted <- ys[accept]
> rejected <- ys[!accept] ; mean(!accept)

[1] 0.3377
```







## Proving the general case

When proving an algorithm generates  $X$  we need to show same cumulative distribution function (recall, we did this for inversion method, transformations, etc.).

Let  $V$  be the random variable produced by the AR algorithm, we need to show that

$$P(V \leq x) = P(X \leq x)$$

for any  $x$  in the support of  $X$ .

## What kind variable is $V$ ?

We know the distribution of  $Y$  and  $U$ , but what can we say about  $V$ ?

Notice that  $V$  is set to  $Y$ , when  $U \leq f(Y)/(cg(Y))$ , a conditional distribution:

$$V = Y \mid U \leq \frac{f(Y)}{c g(Y)}$$

For events  $A$  and  $B$ , recall the definition of conditional probability:

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

## Proof

$$P(V \leq x) = P\left(Y \leq x \mid U \leq \frac{f(Y)}{c g(Y)}\right) = \frac{P\left(Y \leq x, U \leq \frac{f(Y)}{c g(Y)}\right)}{P\left(U \leq \frac{f(Y)}{c g(Y)}\right)}$$

We'll take the numerator and denominator separately.

In both cases, think of writing  $P(\dots) = E(I(\dots))$ , e.g.:

$$P\left(Y \leq x, U \leq \frac{f(Y)}{c g(Y)}\right) = E\left(I\left(Y \leq x, U \leq \frac{f(Y)}{c g(Y)}\right)\right)$$

## Numerator

Notice (a) by independence, the joint density is  $1 \times g(y)$  and (b)  $I(A, B) = I(A)I(B)$ .

$$\begin{aligned}P\left(Y \leq x, U \leq \frac{f(Y)}{c g(Y)}\right) &= \int_{-\infty}^{\infty} \int_0^1 I(y \leq x) I(u \leq f(y)/(c g(y))) g(y) du dy \\&= \int_{-\infty}^{\infty} I(y \leq x) g(y) \left[ \int_0^1 I(u \leq f(y)/(c g(y))) du \right] dy \\&= \int_{-\infty}^x g(y) \left[ \int_0^{f(y)/(c g(y))} 1 du \right] dy \\&= \int_{-\infty}^x g(y) \frac{f(y)}{c g(y)} dy \\&= \frac{1}{c} \int_{-\infty}^x f(y) dy = \frac{1}{c} P(X \leq x)\end{aligned}$$

$$P\left(U \leq \frac{f(Y)}{c g(Y)}\right) = E\left(I\left(U \leq \frac{f(Y)}{c g(Y)}\right)\right)$$

By similar logic,

$$\begin{aligned} P\left(U \leq \frac{f(Y)}{c g(Y)}\right) &= \int_{-\infty}^{\infty} g(y) \left[ \int_0^{f(y)/(cg(y))} 1 \, du \right] dy \\ &= \int_{-\infty}^{\infty} g(y) \frac{f(y)}{cg(y)} dy \\ &= \int_{-\infty}^{\infty} \frac{1}{c} f(y) dy = \frac{1}{c} \end{aligned}$$

Taking the ratio, the  $(1/c)$  cancels, so  $P(V \leq x) = P(X \leq x)$ .

## Example: Truncated Normal

We've used **rejection sampling algorithm** for the Truncated Normal before:

$$Y = Z \mid 0 \leq Z \leq 1, Z \sim N(0, 1)$$

To use AR, we need the **density function**. To get this, start with the CDF:

$$\begin{aligned} F_Y(y) &= P(Y \leq y) = P(Z \leq y \mid 0 \leq Z \leq 1) = \frac{P(Z \leq y, 0 \leq Z \leq 1)}{P(0 \leq Z \leq 1)} \\ &= \frac{1}{P(Z \leq 1) - P(Z \leq 0)} \int_0^y \phi(x) dx \end{aligned}$$

where  $\phi$  is the PDF of the standard Normal distribution.

By definition, the PDF is the derivative of the CDF, so for  $0 \leq y \leq 1$ ,

$$f(y) = \frac{\phi(y)}{P(Z \leq 1) - P(Z \leq 0)}$$

## Implementing in R

```
> (scaling_const <- pnorm(1) - pnorm(0))
```


```
[1] 0.3413
```

Then we compute the density function as:

```
> truncated <- function(x) {  
+   ifelse(x >= 0 & x <= 1,  
+         dnorm(x) / scaling_const,  
+         0)  
+ }
```



## Picking a candidate density: Uniform



## Implementing


```
> ys <- runif(1000)
> const <- truncated(0)
> ratios <- truncated(ys) / (const * 1) #  $g(y) = 1$ 

> us <- runif(1000)
> accept_uniform <- us < ratios
> accepted_uniform <- ys[accept_uniform]
> rejected_uniform <- ys[!accept_uniform]
```

Estimating  $E(X)$ :

```
> mean(accepted_uniform)

[1] 0.4684
```



We rejected 15.1% of the candidate draws. Can we do better?

## Fitting a closer candidate distribution

Consider the density:

$$g(y) = \frac{2}{3}(2 - y), \quad 0 \leq y \leq 1$$


with quantile function

$$Q_y(p) = 2 - \sqrt{4 - 3p}$$

Need to find a  $c$  such that

$$c \times \frac{g(x)}{f(x)} \geq 1 \Rightarrow c \geq \frac{f(x)}{g(x)}, 0 \leq x \leq 1$$


Notice that the **ratio  $f(x)/g(x)$  is increasing**:



Then the maximum ratio is at  $x = 1$ .

$$c = f(1)/g(1) = 0.363 \Rightarrow \frac{cg(x)}{f(x)} \geq 1, 0 \leq x \leq 1$$

```
> g <- function(y) { (2/3) * (2 - y) }  
> qg <- function(p) { 2 - sqrt(4 - 3 * p) }  
> ys <- qg(runif(1000))  
> const <- truncated(1) / g(1)  
> ratios <- truncated(ys) / (const * g(ys))  
  
> us <- runif(1000)  
> accept_g <- us < ratios  
> accepted_g <- ys[accept_g]  
> rejected_g <- ys[!accept_g]
```



Now we reject only 5.1% of the candidate draws.

## Comparing Efficiency

Recall, we define the **efficiency of an estimator** as the **variance of the sampling distribution** for that estimator.

For both the uniform and tuned candidates, the variance is given by:

$$\frac{1}{n} \text{Var}(Y)$$

. The  $\text{Var}(Y)$  term is the same in both cases, so when calculating the **relative efficiency** we only consider the resulting (non-rejected) sample size:

```
> sum(accept_g) / sum(accept_uniform)
```

```
[1] 1.118
```

Interpretation: the method using  $g$  is 10% more efficient than the uniform method.



## A bimodal density

Suppose we have a density of the form:


$$f(x) = 0.25\phi(x + 2) + 0.75\phi(x - 1), -\infty < x < \infty$$

where  $\phi$  is the standard Normal PDF.

```
> f <- function(x) { 0.25 * dnorm(x + 2) + 0.75 * dnorm(x - 1) }
```

A density like this is called **bimodal** because it contains two local maxima.

## Density of $f$



## What candidate distribution?

Our initial thought might be to use  $\phi(x)$  (standard Normal) as the candidate.

The only problem with that is that


$$\lim_{x \rightarrow \pm\infty} f(x)/\phi(x) = \infty$$

(i.e., no  $c$  exists that uniformly bounds  $f$ ).


We saw that the Cauchy distribution has **fat tails**, perhaps that would be useful?

$$g(x) = \frac{1}{\pi(1+x)^2}$$

## Ratio of $f(x)/\phi(x)$




## Ratio of $f(x)/g(x)$



## Lazy mode: find $c$ by evaluation

```
> h <- function(x) { f(x) / dcauchy(x) }  
> xs <- seq(-4, 4, length.out = 1000)  
> (const <- max(h(xs)))  
  
[1] 2.898
```

## Plotting Distributions




## Putting it together

```
> k <- 10000
> ys <- rcauchy(k)
> ratios <- f(ys) / (const * dcauchy(ys))
> us <- runif(k)
> accept <- us < ratios
> x <- ys[accept] ; mean(accept)

[1] 0.3526
```





## What if we got $c$ wrong?

We found  $c = 2.898$ . What if we were wrong? How would the distribution change?  
Too large, too small?

```
> wrong_consts <- c(2 * const,  
+                   const,  
+                   0.75 * const,  
+                   0.25 * const,  
+                   0.1 * const,  
+                   0)
```

## Functions to handle AR

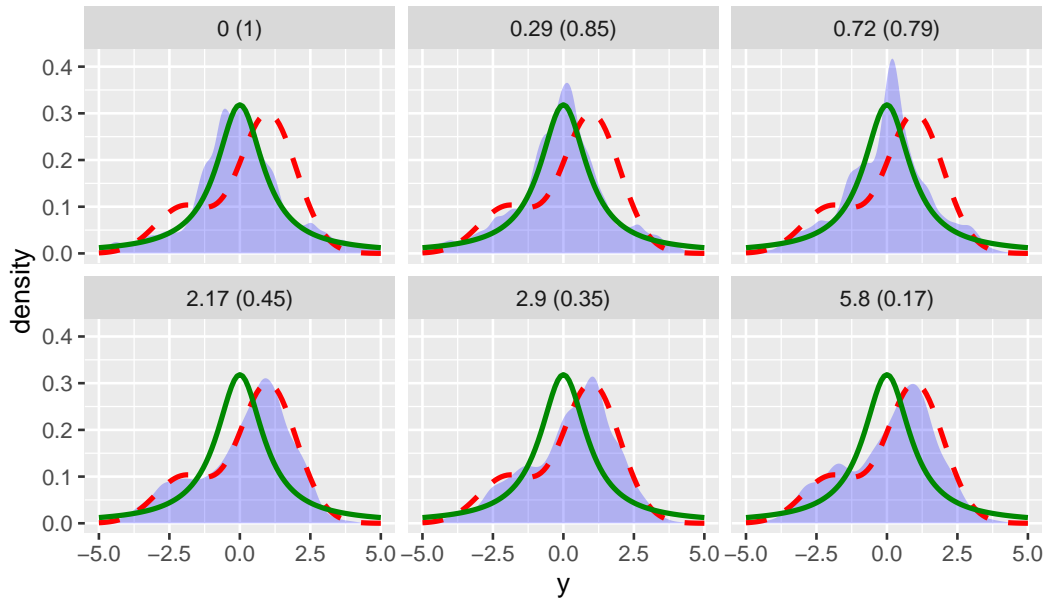
For numerical reasons, we'll rewrite  $U \leq f(Y)/(cg(Y))$  as

$$cg(Y) U \leq f(Y)$$

```
> accept_reject <- function(n, f, g, rg, const) {  
+   ys <- rg(n)  
+   us <- runif(k)  
+   accept <- us * const * g(ys) <= f(ys)  
+   data.frame(y = ys, accept = accept)  
+ }
```

## Using the bad constants

```
> bimodal <- function(const) {  
+   accept_reject(1000, f, dcauchy, rcauchy, const)  
+ }  
  
> wrong <- map_dfr(wrong_consts, .id = "constant", bimodal)  
> wrong$constant <- wrong_consts[as.numeric(wrong$constant)] # get the orig
```



## Accept-Reject Summary

- For a density  $f$ , use a density  $g$  for the sampling.
- Draw uniform values to decide to accept or reject draws for  $g$ .
- Need to figure out the scaling constant  $c$  to make  $cg(y)/f(y) \geq 1$
- Generally, we accept about  $100 \times (1/c)\%$  for the candidates, so making  $c$  small is useful.
- Picking  $c$  too large is ok, too small leads to draws from  $g$ .