

# The Bootstrap

---

Mark M. Fredrickson ([mfredric@umich.edu](mailto:mfredric@umich.edu))

Computational Methods in Statistics and Data Science (Stats 406)

## **Blood Pressure and Low Dose Aspirin**

---

# National Health And Nutrition Examination Survey


The **National Health And Nutrition Examination Survey (NHANES)** provides survey data on the dietary and health habits of people in the United States.

For many years, **low dose aspirin** was thought to be beneficial for those at risk of **heart disease**.

Variables:

- Respondent's self-reported use of low dose aspirin
- Diastolic and systolic blood pressure

See additional slides at end of lecture for data loading.



# The Bootstrap

---

## Set up

Suppose we have a sample of **size**  $n$ , which we will **assume are independent, identically distributed**.

$$X_i \sim F, \quad i = 1, \dots, n, \quad (\text{independent})$$

We want to estimate  $\theta = E(h(X))$  using a **statistic**  $T(X_1, \dots, X_n) = T$  and construct a **confidence interval** for  $\theta$ .

## Usual Methods

We need to know the **sampling distribution of  $T$**  (the distribution of  $T$  if we picked many samples of size  $n$ ).

- **Derive the distribution** of  $T$  from **assumptions** about the  $X_i$ .
- Find a **large sample approximation** for  $T$  and assume we have enough data for it to apply.
- **Fix a specific distribution** and use **Monte Carlo**.

Alternative: What if we could estimate the **distribution of  $X$**  and then use the **inversion method** to draw from our estimate?

## Estimating $F$

Recall the definition of  $F(x)$ :

$$F(x) = \Pr(X \leq x) = E(I(X \leq x))$$

We developed tools to estimate things like  $E(I(X \leq x))$ :

$$\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n I(X_i \leq x)$$

This is the **empirical cumulative distribution function**.



## Inversion method

Recall, the **inversion method** draws  $U \sim U(0, 1)$  and then generates  $X = Q_X(U)$ . Here, we will use the **empirical quantile function**.

Notice that  $\hat{F}_n(x)$  defines a **discrete distribution**, so the **empirical quantile function** is:

$$\hat{Q}_n(u) = x \text{ such that } \hat{F}_n(x - \epsilon) < u \leq \hat{F}_n(x), \forall \epsilon > 0$$

Notice: the only possible  $x$  values will be those that appear in the sample.

Claim: this is equivalent to picking an observation **uniformly at random**:

$$P(\hat{Q}(U) = X_i) = \frac{1}{n}, \quad \forall i$$

## Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Observe:  $\sum I(X_i \leq X_{(j)}) = j$

$$\begin{aligned} P(\hat{Q}(U) = X_{(j)}) &= P\left(\hat{F}(X_{(j-1)}) < U \leq \hat{F}(X_{(j)})\right) \\ &= P\left(\frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j-1)}) < U \leq \frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j)})\right) \\ &= P((j-1)/n < U \leq j/n) \\ &= P(U \leq j/n) - P(U \leq (j-1)/n) \\ &= \frac{j}{n} - \frac{j-1}{n} = \frac{1}{n} \end{aligned}$$

This process is equivalent to picking one of the  $X_i$  uniformly at random.

# The Bootstrap

Key idea: estimate the **sampling distribution of  $T$**  by drawing many samples from the **estimated distribution for  $X$** .


There are  $n^n$  possible samples, so we can't enumerate them all. So we will take a sample (**the bootstrap sample**).

Let  $X_1^*, X_2^*, \dots, X_n^*$  be a sample picked from the original  $n$  observations, taken with replacement. We will do this  $B$  times.

For each sample, compute  $T(X_1^*, X_2^*, \dots, X_n^*) = T^*$ .

## NHANES systolic measurements

Here's the distribution of systolic blood pressure (average) readings for the NHANES survey:



Let's estimate the **population mean of systolic blood pressure** (assuming NHANES is IID from US pop).

## The trimmed mean

The previous distribution displayed a mild amount of **right-skew**.

While the CLT holds for large enough  $n$ , the **trimmed mean** might be a better choice.

Recall,  $p \in [0, 1]$ :

- Discard the lower  $p/2$  and upper  $p/2$  proportion of the data
- Compute the mean on the remaining observations

What is the **sampling distribution** of the trimmed mean for our data?


## Trimmed Mean

```
> trimmed_mean <- function(x, p) {  
+   xqs <- quantile(x, c(p/2, 1 - (p/2)))  
+   keep <- x > xqs[1] & x < xqs[2]  
+   mean(x[keep])  
+ }  
  
> (observed_mean <- mean(sys_mean))  
  
[1] 127.6  
  
> (observed_trim <- trimmed_mean(sys_mean, p = 0.2))  
  
[1] 126.2
```

## Bootstrap Sample

```
> B <- 10000  
> n <- length(sys_mean)  
> bootstrap_samples <- rerun(B,  
+   sample(sys_mean, size = n, replace = TRUE))  
  
> bootstrap_means <- map_dbl(bootstrap_samples, mean)  
> bootstrap_trims <- map_dbl(bootstrap_samples, trimmed_mean, p = 0.2)
```

## Comparing Sampling Distributions





## Using the distribution

We are **estimating the sampling distribution** for these statistics. What information might we want to know about?

- Variance
- Quantiles
- Bias (would need to know true parameter) and MSE

We'll use the first two in order to **develop confidence intervals** and return to bias and MSE later.

# Bootstrap Confidence Intervals

---

## Bootstrap confidence intervals

Our goal is to estimate some  $\theta$  using  $\hat{\theta} = T(X_1, X_2, \dots, X_n)$  (we get this without the bootstrap).

We use the bootstrap to understand the **sampling distribution** of  $\hat{\theta}$ . We can use this to create **confidence intervals** in three basic ways:

- **Large sample approximation** intervals  $\hat{\theta} \pm z_{\alpha/2} \hat{\sigma}_{\hat{\theta}}$
- Combining the **estimator on the original sample** with the **bootstrap distribution** in two different ways.

## Large sample (Normal approximation)

Many statistics have the property that in **large samples**,

$$T \sim N(\theta, \sigma_T^2) \quad (\text{approx.})$$

Using **test inversion** we could create a confidence interval as:

$$T \pm z_{\alpha/2} \sigma_T$$

Even if this is known, it may be difficult to compute or estimate  $\sigma_T$ . We substitute the **standard deviation of the bootstrap distribution**:

$$T \pm z_{\alpha} \sqrt{\frac{1}{B-1} \sum_{i=j}^B (T_j^* - \bar{T}^*)^2}$$

## Bootstrap intervals vs. usual $t$ -intervals

We derived the following interval:

$$T \pm z_{\alpha} \sqrt{\frac{1}{B-1} \sum_{i=j}^B (T_j^* - \bar{T}^*)^2}$$

Notice that this interval differs from the usual  $t$  based interval in three ways:

- We used a Normal quantile instead of a  $t$ -quantile. We could have used a  $t$ -quantile, but with large  $B$  it is basically equivalent.
- We used  $T$  instead of  $\bar{T}^*$  as the center of the interval (observed estimator instead of mean of the bootstrap distribution).
- We are estimating the variance of  $T$  directly, so there is no  $1/\sqrt{n}$  term (variance of sample mean is  $\text{Var}(\bar{X}) = (1/n)\text{Var}(X)$ ).

## Normal interval for systolic BP

The trimmed mean is a good example of a statistic that has an asymptotic Normal distribution (see Stigler (1973)) but a tricky variance.

```
> observed_trim + c(1, -1) * qnorm(0.025) * sd(bootstrap_trims)
```

```
[1] 125.6 126.9
```

## Basic intervals

As before let  $T = T(X_1, X_2, \dots, X_n)$  (the **observed statistic**) and  $T^* = T(X_1^*, X_2^*, \dots, X_n^*)$  (the **bootstrap statistic**).

Suppose we want to create intervals of the form,  $[T - a, T + b]$  so that

$$P(T - a < \theta < T + b) \geq 1 - \alpha$$

(i.e.,  $[T - a, T + b]$  is a valid CI).

Rewriting, we have the equivalent formulation:

$$P(-b \leq T - \theta \leq a) \geq 1 - \alpha$$

To pick  $-b$  and  $a$ , we need the distribution of  $T - \theta$ .

## Estimating $T - \theta$ using the bootstrap

The main goal of bootstrapping is **approximate the distribution of  $T$  using  $T^*$** . So that

$$P(-b \leq T - \theta \leq a) \approx P(-b \leq T^* - \theta \leq a)$$

We know the distribution of  $T^*$  (more properly, have an MC estimate of it), **what about  $\theta$** ?

Provided  $T$  is a good estimator of  $\theta$  (i.e.,  $T \xrightarrow{P} \theta$ ), then

$$P(-b \leq T^* - \theta \leq a) \approx P(-b \leq T^* - T \leq a)$$

This implies that we can pick  $-b = T_{\alpha/2}^* - T$  and  $a = T_{1-\alpha/2}^* - T$ .



## Putting it all together

We wanted CIs of the form:

$$P(\theta \in [T - a, T + b]) \geq 1 - \alpha$$

We found  $a = T_{1-\alpha/2}^* - T$  and  $b = T - T_{\alpha/2}^*$ .

$$[T - a, T + b] = [2T - T_{1-\alpha/2}^*, 2T - T_{\alpha/2}^*]$$

We call these **basic bootstrap confidence intervals**.

## Applying to previous examples

```
> ba_trim <- quantile(bootstrap_trims, c(0.975, 0.025))  
> basic_trim <- 2 * observed_trim - ba_trim  
> names(basic_trim) <- NULL # quantile adds some names, we don't need them  
> basic_trim  
  
[1] 125.7 126.9
```

## Percentile Intervals

During the previous approach, we were looking for  $a$  and  $b$  such that

$$P(-b \leq T - \theta \leq a) \geq 1 - \alpha$$

and we had the approximation:

$$P(T_{\alpha/2}^* - T \leq T - \theta \leq T_{1-\alpha/2}^* - T) = P(T_{\alpha/2}^* \leq 2T - \theta \leq T_{1-\alpha/2}^*)$$

Again, since  $T \xrightarrow{P} \theta$ ,  $2T \approx 2\theta$ , so

$$P(T_{\alpha/2}^* \leq 2T - \theta \leq T_{1-\alpha/2}^*) \approx P(T_{\alpha/2}^* \leq 2\theta - \theta \leq T_{1-\alpha/2}^*) = P(T_{\alpha/2}^* \leq \theta \leq T_{1-\alpha/2}^*)$$

which suggests intervals:

$$[T_{\alpha/2}^*, T_{1-\alpha/2}^*]$$

These **percentile intervals** also often form the basis for other refinements.

## Trimmed mean example

```
> quantile(bootstrap_trims, c(0.025, 0.975))
```

```
2.5% 97.5%
```

```
125.6 126.8
```

## Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.
- If you are going to compute the full bootstrap, **why use the Normal approximation?**
- The basic and percentile intervals will always have **the same width**.
- All methods are **approximations** based on the assumptions that:
  - $T \approx \theta$  ( $n$  is large)
  - $\hat{F} \approx F$  ( $n$  is large)
  - $F^* \approx F$  ( $B$  is large)

## Using R's `boot` package

Much of the code for running bootstrap estimates is repetitive. Let's use the `boot` package instead.

We need to set up our test statistic function to take a copy of the data and an index for the particular bootstrap sample.

```
> trimmed_mean_boot <- function(x, index, p = 0.1) {  
+   trimmed_mean(x[index], p = p)  
+ }  
> library(boot)  
> boot_tm <- boot(sys_mean, statistic = trimmed_mean_boot, p = 0.2, R = B)
```

```
> boot.ci(boot_tm, type = c("norm", "basic", "perc"))
```

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 10000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_tm, type = c("norm", "basic", "perc"))
```


Intervals :

Level	Normal	Basic	Percentile
95%	(125.7, 126.9 )	(125.7, 127.0 )	(125.5, 126.8 )

Calculations and Intervals on Original Scale

# Estimating the correlation between systolic and diastolic BP

Research question: Are systolic and diastolic pressure linearly related?






## Bootstrapping the Correlation Coefficient

```
> cor_boot_stat <- function(x, index) {  
+   cor(x[index, 1], x[index, 2], use = "complete")  
+ }  
> library(boot)  
> boot_cor <- boot(nhanes[, c("sys_mean", "dia_mean")],  
+   statistic = cor_boot_stat,  
+   R = 1000)
```

## Correlation coefficient distribution



```
> boot.ci(boot_cor, type = c("norm", "basic", "perc"))
```

# BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_cor, type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	( 0.2269, 0.3117 )	( 0.2266, 0.3118 )	( 0.2262, 0.3114 )

Calculations and Intervals on Original Scale

Suppose  $X_1, \dots, X_n \sim N(\mu, \sigma^2)$  (independent). Then

$$W = \frac{\bar{X} - \mu}{(S^2/n)^{1/2}}$$

has a Student's  $t$ -distribution with  $n - 1$  degrees of freedom.

More generally we say that a statistic is **studentized** if we subtract off a hypothesized location parameter and divide by an estimate of the standard deviation of the estimator.

## Bootstrap-t (percentile) confidence intervals

We noted that **percentile** confidence intervals are frequently the basis of **improved confidence intervals**.

Define the “studentized” bootstrap replicate

$$W^* = \frac{T^* - T}{\hat{\sigma}^*}$$

Undo the studentization to get back to the  $T$  scale:

$$[T - \hat{\sigma} W_{1-\alpha/2}^*, T - \hat{\sigma} W_{\alpha/2}^*]$$

## Variance Estimators


In the previous algorithm, we used **two different variance estimators** (for notational ease, I'm going to write these using standard deviations instead):

- $\hat{\sigma}^*$ : estimates  $\text{Var}(T^*)^{1/2}$  based on a particular bootstrap sample
- $\hat{\sigma}$ : estimates  $\text{Var}(T)^{1/2}$  based on the original sample

For either of these we could use

- Analytical estimator (e.g.,  $\text{Var}(\bar{X}) = (1/n)\text{Var}(X)$  and estimate  $\text{Var}(X)$  using sample variance statistic  $S_x^2$ )
- Bootstrap estimate of variance (“nested bootstrap”)
- The Jackknife (which we'll discuss a bit later)

## Log Ratio of Systolic to Diastolic



## Bootstrapping the mean

```
> library(boot)
> mean_boot <- function(x, index) { mean(x[index]) }
> boot_mean <- boot(log(sysdia_ratio), statistic = mean_boot, R = 1000)
```



```
> boot.ci(boot_mean, type = c("norm", "basic", "perc"))
```

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_mean, type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	( 0.5931, 0.6086 )	( 0.5929, 0.6085 )	( 0.5935, 0.6090 )

Calculations and Intervals on Original Scale

## Bootstrap-t: sample variance estimator

```
> B <- 1000  
> lsdr <- log(sysdia_ratio)  
> n <- length(lsdr)  
> est_t <- mean(lsdr)  
> est_var_t <- var(lsdr) / n
```

```
> boot_sv <- replicate(B, {  
+   xstar <- sample(lsdrr, replace = TRUE)  
+   (mean(xstar) - est_t) / sqrt(var(xstar) / n)  
+ })  
> (boot_ci_svar <- est_t - sqrt(est_var_t) *  
+   quantile(boot_sv, c(0.975, 0.025)))  
  
      97.5%      2.5%  
0.5932055 0.6088884
```

## Bootstrap-t: Nested bootstrap

```
> boot_for_var <- replicate(100, {  
+   xstar <- sample(lsd, replace = TRUE)  
+   mean(xstar)  
+ })  
> boot_var_est <- var(boot_for_var)
```

```

> boot_boot <- replicate(100, {
+   xstar <- sample(lsdrr, replace = TRUE)
+   xstar_boot <- replicate(100, {
+     xstarstar <- sample(xstar, replace = TRUE)
+     mean(xstarstar)
+   })
+   (mean(xstar) - est_t) / sd(xstar_boot)
+ })
> (boot_ci_boot <- est_t - sqrt(boot_var_est) *
+   quantile(boot_boot, c(0.975, 0.025)))

```

```

      97.5%      2.5%
0.5942140 0.6081809

```

## Using the boot package

If we return **two values**, the boot package will treat the first as  $T^*$  and the second as  $\hat{\sigma}_*^2$ .

```
> mean_var <- function(x, index) {  
+   xstar <- x[index]  
+   n <- length(xstar)  
+   c(mean(xstar), var(xstar) / n)  
+ }  
> boot_both <- boot(lsd, mean_var, R = 1000)
```

```
> boot.ci(boot_both, type = 'stud')
```

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_both, type = "stud")
```

Intervals :

Level	Studentized
-------	-------------

95%	( 0.5927, 0.6089 )
-----	--------------------

Calculations and Intervals on Original Scale

## Comparing CIs

	Low		High	Rel. Width
Basic	0.592299625	0.608321349		1.000000000
Percentile	0.593614827	0.609636551		1.000000000
Studentized	0.592678187	0.608857458		1.009833294



## Nested bootstrap of the median

With long tailed data (like the log-ratio were using), **the median** may be a better measure of **central tendency** than the mean.

We can't use sample variance estimate for the median, so we'll use **nested bootstrap**.

```
> median_idx <- function(data, idx) {  
+   median(data[idx])  
+ }  
  
> median_nested <- function(data, idx) {  
+   xstar <- data[idx]  
+   meds <- boot(xstar, median_idx, R = 100)$t # just the T values  
+   return(c(median(xstar), var(meds)))  
+ }
```

## Bootstrapping with Parallel Library

```
> library(parallel)
> cl <- makeCluster(detectCores())
> ## load the nested bootstrap components on the cluster
> ignore <- clusterEvalQ(cl, library(boot))
> clusterExport(cl, c("median_idx", "median_nested"))
> boot_median <- boot(lsdrr, median_nested, R = 1000,
+   parallel = "snow", cl = cl, ncpus = detectCores())
> stopCluster(cl)
```

## Confidence Intervals

```
> boot.ci(boot_median, type = "stud")
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_median, type = "stud")
```

Intervals :

Level	Studentized
-------	-------------

95%	( 0.5423, 0.5550 )
-----	--------------------

Calculations and Intervals on Original Scale

## Assessing CI Coverage

---

# Confidence Coefficient

For a **parameter**  $\theta$  and the **random interval**  $[A, B]$  we define the **confidence coefficient**  $c$  as

$$c = P(A \leq \theta, B \geq \theta)$$

If  $c \geq 1 - \alpha$  then  $[A, B]$  is a valid  $(1 - \alpha) \times 100\%$  **Confidence Interval**.

It can be the case that a procedure we claim is a valid CI has  $c < 1 - \alpha$ . We need to evaluate our procedures to make sure this doesn't happen.

# Large Sample Intervals

For

$$X \sim F, E(X) = \mu, \text{Var}(X) < \infty$$

the **central limit theorem** suggests that  $\bar{X}$  is approximately

$$\bar{X} \sim N(\mu, \text{Var}(X)/n)$$

This suggests that the usual  $t$ -intervals are **approximate confidence intervals**:

$$P(\bar{X} - t_{1-\alpha/2}s/\sqrt{n} \leq \mu \leq \bar{X} + t_{1-\alpha/2}s/\sqrt{n}) \approx 1 - \alpha$$

## Quality of Approximation

As we saw in HW3, the **quality of this approximation depends on  $F$** , the distribution of  $X$ .

We investigated the following when  $\mu = 1/2$  and found

- $n = 20, X_i \sim \text{Laplace}(1/2)$ : good approximation
- $n = 20, X_i \sim \text{Exp}(2)$ : poor approximation
- $n = 500, X_i \sim \text{Exp}(2)$ : good approximation

Difficult problem: in general, is the Normal approximation valid?

## Quick review of HW3 Results: Laplace

```
> k <- 10000
> laplace_intervals <- rerun(k, {
+   t.test(rlaplace(20, mean = 1/2), conf.level = 0.95)$conf.int
+ })
> a1 <- map_dbl(laplace_intervals, ~ .x[1] <= 1/2 & .x[2] >= 1/2)
> binom.test(sum(a1), k, conf.level = 0.99)$conf.int

[1] 0.9439078 0.9552592
attr(,"conf.level")
[1] 0.99
```



## Quick review of HW3 Results: Exponential

```
> exponential_intervals <- rerun(k, {  
+   t.test(rexp(20, rate = 2), conf.level = 0.95)$conf.int  
+ })  
> a2 <- map_dbl(exponential_intervals, ~ .x[1] <= 1/2 & .x[2] >= 1/2)  
> binom.test(sum(a2), k, conf.level = 0.99)$conf.int  
  
[1] 0.9132797 0.9273164  
attr(,"conf.level")  
[1] 0.99
```

## Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

- Generate sample  $n$  from known distribution
- Perform bootstrap procedure  $B$  times
- Compute Normal theory, basic and percentile intervals
- Repeat  $k$  times

Note, requires about  $n \times B \times k$  steps. Let's **parallelize!**

## Setting up local cluster

```
> library(boot)
> library(parallel)
> cl <- makeCluster(detectCores() - 1)
> ## Example usage:
> mean_boot <- function(x, idx) { mean(x[idx]) }
> a <- boot(rlaplace(20, 1/2), mean_boot, R = 1000,
+           parallel = "snow",
+           cl = cl,
+           ncpus = 2) %>% # bug: must be greater than 1
+           boot.ci(type = c("norm", "basic", "perc"))
```

## Formatted results

```
> print(a)
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = ., type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	( 0.3568, 1.1857 )	( 0.3439, 1.1986 )	( 0.3674, 1.2221 )

Calculations and Intervals on Original Scale

## What is the structure of `boot.ci` result?

```
> names(a)
```

```
[1] "R"          "t0"          "call"         "normal"      "basic"       "percent"
```

```
> a$normal ; a$basic ; a$percent
```

```
      conf
```

```
[1,] 0.95 0.3567675 1.185686
```

```
      conf
```

```
[1,] 0.95 975.98 25.03 0.3439484 1.198577
```

```
      conf
```

```
[1,] 0.95 25.03 975.98 0.367446 1.222074
```

## Pulling out just the intervals

```
> getCIs <- function(boot_ci_result) {  
+   with(boot_ci_result,  
+       matrix(c(normal[2:3],  
+               basic[4:5],  
+               percent[4:5]), nrow = 2))  
+ }  
> getCIs(a)
```

	[,1]	[,2]	[,3]
[1,]	0.3567675	0.3439484	0.367446
[2,]	1.1856857	1.1985765	1.222074

## Putting it together: Laplace

```
> k <- 1000 ; R <- 1000
> laplace_bootstrap_cis <- rerun(k, {
+   boot(rlaplace(20, 1/2), mean_boot, R = R,
+       parallel = "snow", cl = cl, ncpus = 2) %>%
+   boot.ci(type = c("norm", "basic", "perc")) %>%
+   getCIs
+ })
```

## Putting it together: Exponential

```
> exp_bootstrap_cis <- rerun(k, {  
+   boot(rexp(20, 2), mean_boot, R = R,  
+       parallel = "snow", cl = cl, ncpus = 2) %>%  
+       boot.ci(type = c("norm", "basic", "perc")) %>%  
+       getCIs  
+ })
```



## Counting Covering CIs

```
> covers <- function(x) { x[1, ] <= 1/2 & x[2, ] >= 1/2 }  
> laplace_bootstrap_covers <- map(laplace_bootstrap_cis, covers) %>%  
+   simplify2array %>% rowSums  
> exp_bootstrap_covers <- map(exp_bootstrap_cis, covers) %>%  
+   simplify2array %>% rowSums
```

## Coverage Rates: Laplace

```
> coverage_ci <- function(x) { binom.test(x, k, conf.level = 0.99)$conf.int  
> map(laplace_bootstrap_covers, coverage_ci)
```

```
[[1]]
```

```
[1] 0.9099948 0.9517603
```

```
[[2]]
```

```
[1] 0.9237124 0.9619124
```

```
[[3]]
```

```
[1] 0.8975841 0.9422873
```

## Coverage Rates: Exponential

```
> map(exp_bootstrap_covers, coverage_ci)
```

```
[[1]]
```

```
[1] 0.8621267 0.9140733
```

```
[[2]]
```

```
[1] 0.8511989 0.9051019
```

```
[[3]]
```

```
[1] 0.8698147 0.9203143
```

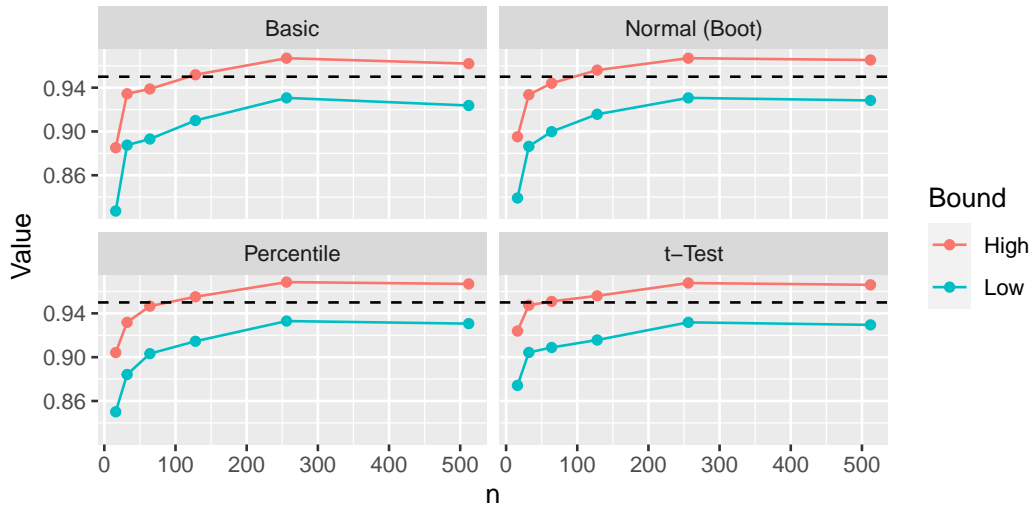
## Why poor coverage rates?

Remember, we need  $\hat{F}(x)$  **to be close to**  $F(x)$  for all  $x$ .

We had a sample size of 20. Is this enough for a good approximation? (Apparently not.)

Let's repeat this procedure at **increasingly large sample sizes** (and compare to `t.test` along the way).

## Plotting results



## Interpreting Coverage Rates

In the previous examples, we saw that **all** of the methods had **confidence coefficients** below the targeted level.

For the t-test based intervals, we know this is because **the CLT Normal approximation only holds as  $n \rightarrow \infty$** .

For the other methods, we only have that  $\hat{F}(x) \xrightarrow{P} F$ , another **law of large numbers** type result. The bootstrap requires the approximation  $\hat{F}(x) \approx F$  to hold, which for small samples is tenuous.

Remember: the primary benefit of the bootstrap is **trading Monte Carlo methods for analytic methods** not magically making new data.

# Summary

- Goal: **estimate**  $\theta$  using an **estimator**  $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$ .
- Need to know the **sampling distribution of  $\hat{\theta}$**  to find **confidence intervals**.
- Main idea: If  $X_i \stackrel{\text{iid}}{\sim} F$ , **estimate**  $\hat{F}$  in order to apply **inversion method**:  
 $\hat{Q}(U) \sim F$  (approximately).
- The **bootstrap sample** of  $X_1^*, X_2^*, \dots, X_n^*$  is used to compute  
 $T^* = T(X_1^*, \dots, X_n^*)$ .
- Since  $T^* \sim T$  (approx.) we can create confidence intervals:
  - **Large sample**:  $T \pm z_{\alpha/2} \sigma_{T^*}$
  - **Basic**:  $[2T - T_{1-\alpha/2}^*, 2T - T_{\alpha/2}^*]$
  - **Percentile**:  $[T_{\alpha/2}^*, T_{1-\alpha/2}^*]$
  - **Studentized**:  $[T - \sigma W_{1-\alpha/2}^*, T - \sigma W_{\alpha/2}^*]$ , where  $W^* = (T^* - T)/\hat{\sigma}^*$  (requires within bootstrap sample variance estimate)