

The Bootstrap

Mark M. Fredrickson (mfredric@umich.edu)

Computational Methods in Statistics and Data Science (Stats 406)

Blood Pressure and Low Dose Aspirin

National Health And Nutrition Examination Survey

The **National Health And Nutrition Examination Survey (NHANES)** provides survey data on the dietary and health habits of people in the United States.

National Health And Nutrition Examination Survey

The **National Health And Nutrition Examination Survey (NHANES)** provides survey data on the dietary and health habits of people in the United States.

For many years, **low dose aspirin** was thought to be beneficial for those at risk of **heart disease**.

National Health And Nutrition Examination Survey


The **National Health And Nutrition Examination Survey (NHANES)** provides survey data on the dietary and health habits of people in the United States.

For many years, **low dose aspirin** was thought to be beneficial for those at risk of **heart disease**.

Variables:

- Respondent's self-reported use of low dose aspirin
- Diastolic and systolic blood pressure

See additional slides at end of lecture for data loading.



The Bootstrap

Set up

Suppose we have a sample of size n , which we will assume are independent, identically distributed.

$$X_i \sim F, \quad i = 1, \dots, n, \quad (\text{independent})$$

Set up

Suppose we have a sample of **size** n , which we will **assume are independent, identically distributed**.

$$X_i \sim F, \quad i = 1, \dots, n, \quad (\text{independent})$$

We want to estimate $\theta = E(h(X))$ using a **statistic** $T(X_1, \dots, X_n) = T$ and construct a **confidence interval** for θ .

We need to know the **sampling distribution of T** (the distribution of T if we picked many samples of size n).

Usual Methods

We need to know the **sampling distribution of T** (the distribution of T if we picked many samples of size n).

- **Derive the distribution** of T from **assumptions** about the X_i .

Usual Methods

We need to know the **sampling distribution of T** (the distribution of T if we picked many samples of size n).

- **Derive the distribution** of T from **assumptions** about the X_i .
- Find a **large sample approximation** for T and assume we have enough data for it to apply.

Usual Methods

We need to know the **sampling distribution of T** (the distribution of T if we picked many samples of size n).

- **Derive the distribution** of T from **assumptions** about the X_i .
- Find a **large sample approximation** for T and assume we have enough data for it to apply.
- **Fix a specific distribution** and use **Monte Carlo**.

Usual Methods

We need to know the **sampling distribution of T** (the distribution of T if we picked many samples of size n).

- **Derive the distribution** of T from **assumptions** about the X_i .
- Find a **large sample approximation** for T and assume we have enough data for it to apply.
- **Fix a specific distribution** and use **Monte Carlo**.

Usual Methods

We need to know the **sampling distribution of T** (the distribution of T if we picked many samples of size n).

- **Derive the distribution** of T from **assumptions** about the X_i .
- Find a **large sample approximation** for T and assume we have enough data for it to apply.
- **Fix a specific distribution** and use **Monte Carlo**.

Alternative: What if we could estimate the **distribution of X** and then use the **inversion method** to draw from our estimate?

Estimating F

Recall the definition of $F(x)$:

$$F(x) = \Pr(X \leq x)$$

Recall the definition of $F(x)$:

$$F(x) = \Pr(X \leq x) = E(I(X \leq x))$$

Estimating F

Recall the definition of $F(x)$:

$$F(x) = \Pr(X \leq x) = E(I(X \leq x))$$

We developed tools to estimate things like $E(I(X \leq x))$:

$$\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n I(X_i \leq x)$$

This is the **empirical cumulative distribution function**.

Inversion method

Recall, the **inversion method** draws $U \sim U(0, 1)$ and then generates $X = Q_X(U)$.
Here, we will use the **empirical quantile function**.

Inversion method

Recall, the **inversion method** draws $U \sim U(0, 1)$ and then generates $X = Q_X(U)$. Here, we will use the **empirical quantile function**.

Notice that $\hat{F}_n(x)$ defines a **discrete distribution**, so the **empirical quantile function** is:

$$\hat{Q}_n(u) = x \text{ such that } \hat{F}_n(x - \epsilon) < u \leq \hat{F}_n(x), \forall \epsilon > 0$$

Inversion method

Recall, the **inversion method** draws $U \sim U(0, 1)$ and then generates $X = Q_X(U)$. Here, we will use the **empirical quantile function**.

Notice that $\hat{F}_n(x)$ defines a **discrete distribution**, so the **empirical quantile function** is:

$$\hat{Q}_n(u) = x \text{ such that } \hat{F}_n(x - \epsilon) < u \leq \hat{F}_n(x), \forall \epsilon > 0$$

Notice: the only possible x values will be those that appear in the sample.

Inversion method

Recall, the **inversion method** draws $U \sim U(0, 1)$ and then generates $X = Q_X(U)$. Here, we will use the **empirical quantile function**.

Notice that $\hat{F}_n(x)$ defines a **discrete distribution**, so the **empirical quantile function** is:

$$\hat{Q}_n(u) = x \text{ such that } \hat{F}_n(x - \epsilon) < u \leq \hat{F}_n(x), \forall \epsilon > 0$$

Notice: the only possible x values will be those that appear in the sample.

Claim: this is equivalent to picking an observation **uniformly at random**:

$$P(\hat{Q}(U) = X_i) = \frac{1}{n}, \quad \forall i$$

Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Observe: $\sum I(X_i \leq X_{(j)}) = j$

Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Observe: $\sum I(X_i \leq X_{(j)}) = j$

$$P(\hat{Q}(U) = X_{(j)}) = P\left(\hat{F}(X_{(j-1)}) < U \leq \hat{F}(X_{(j)})\right)$$

Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Observe: $\sum I(X_i \leq X_{(j)}) = j$

$$\begin{aligned} P(\hat{Q}(U) = X_{(j)}) &= P\left(\hat{F}(X_{(j-1)}) < U \leq \hat{F}(X_{(j)})\right) \\ &= P\left(\frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j-1)}) < U \leq \frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j)})\right) \end{aligned}$$

Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Observe: $\sum I(X_i \leq X_{(j)}) = j$

$$\begin{aligned} P(\hat{Q}(U) = X_{(j)}) &= P\left(\hat{F}(X_{(j-1)}) < U \leq \hat{F}(X_{(j)})\right) \\ &= P\left(\frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j-1)}) < U \leq \frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j)})\right) \\ &= P((j-1)/n < U \leq j/n) \end{aligned}$$

Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Observe: $\sum I(X_i \leq X_{(j)}) = j$

$$\begin{aligned} P(\hat{Q}(U) = X_{(j)}) &= P\left(\hat{F}(X_{(j-1)}) < U \leq \hat{F}(X_{(j)})\right) \\ &= P\left(\frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j-1)}) < U \leq \frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j)})\right) \\ &= P((j-1)/n < U \leq j/n) \\ &= P(U \leq j/n) - P(U \leq (j-1)/n) \end{aligned}$$

Proof

Breaking ties arbitrarily, relabel our data such that

$$X_{(1)} \leq X_{(2)} \leq \cdots \leq X_{(n)}$$

Observe: $\sum I(X_i \leq X_{(j)}) = j$

$$\begin{aligned} P(\hat{Q}(U) = X_{(j)}) &= P\left(\hat{F}(X_{(j-1)}) < U \leq \hat{F}(X_{(j)})\right) \\ &= P\left(\frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j-1)}) < U \leq \frac{1}{n} \sum_{i=1}^n I(X_i \leq X_{(j)})\right) \\ &= P((j-1)/n < U \leq j/n) \\ &= P(U \leq j/n) - P(U \leq (j-1)/n) \\ &= \frac{j}{n} - \frac{j-1}{n} = \frac{1}{n} \end{aligned}$$

This process is equivalent to picking one of the X_i uniformly at random.

The Bootstrap

Key idea: estimate the **sampling distribution of T** by drawing many samples from the **estimated distribution for X** .

The Bootstrap

Key idea: estimate the **sampling distribution of T** by drawing many samples from the **estimated distribution for X** .

There are n^n possible samples, so we can't enumerate them all. So we will take a sample (**the bootstrap sample**).

The Bootstrap

Key idea: estimate the **sampling distribution of T** by drawing many samples from the **estimated distribution for X** .

There are n^n possible samples, so we can't enumerate them all. So we will take a sample (**the bootstrap sample**).

Let $X_1^*, X_2^*, \dots, X_n^*$ be a sample picked from the original n observations, taken with replacement. We will do this B times.

The Bootstrap

Key idea: estimate the **sampling distribution of T** by drawing many samples from the **estimated distribution for X** .


There are n^n possible samples, so we can't enumerate them all. So we will take a sample (**the bootstrap sample**).

Let $X_1^*, X_2^*, \dots, X_n^*$ be a sample picked from the original n observations, taken with replacement. We will do this B times.

For each sample, compute $T(X_1^*, X_2^*, \dots, X_n^*) = T^*$.


NHANES systolic measurements

Here's the distribution of systolic blood pressure (average) readings for the NHANES survey:



NHANES systolic measurements

Here's the distribution of systolic blood pressure (average) readings for the NHANES survey:



Let's estimate the **population mean of systolic blood pressure** (assuming NHANES is IID from US pop).

The trimmed mean

The previous distribution displayed a mild amount of **right-skew**.

The trimmed mean

The previous distribution displayed a mild amount of **right-skew**.

While the CLT holds for large enough n , the **trimmed mean** might be a better choice.

The trimmed mean

The previous distribution displayed a mild amount of **right-skew**.

While the CLT holds for large enough n , the **trimmed mean** might be a better choice.

Recall, $p \in [0, 1]$:

- Discard the lower $p/2$ and upper $p/2$ proportion of the data
- Compute the mean on the remaining observations

The trimmed mean

The previous distribution displayed a mild amount of **right-skew**.

While the CLT holds for large enough n , the **trimmed mean** might be a better choice.

Recall, $p \in [0, 1]$:

- Discard the lower $p/2$ and upper $p/2$ proportion of the data
- Compute the mean on the remaining observations

What is the **sampling distribution** of the trimmed mean for our data?

Trimmed Mean

```
> trimmed_mean <- function(x, p) {  
+   xqs <- quantile(x, c(p/2, 1 - (p/2)))  
+   keep <- x > xqs[1] & x < xqs[2]  
+   mean(x[keep])  
+ }  
  
> (observed_mean <- mean(sys_mean))  
  
[1] 127.6  
  
> (observed_trim <- trimmed_mean(sys_mean, p = 0.2))  
  
[1] 126.2
```


Bootstrap Sample


```
> B <- 10000  
> n <- length(sys_mean)  
> bootstrap_samples <- rerun(B,  
+   sample(sys_mean, size = n, replace = TRUE))
```

Bootstrap Sample

```
> B <- 10000
> n <- length(sys_mean)
> bootstrap_samples <- rerun(B,
+   sample(sys_mean, size = n, replace = TRUE))

> bootstrap_means <- map_dbl(bootstrap_samples, mean)
> bootstrap_trims <- map_dbl(bootstrap_samples, trimmed_mean, p = 0.2)
```

Comparing Sampling Distributions



Using the distribution

We are **estimating the sampling distribution** for these statistics. What information might we want to know about?

Using the distribution

We are **estimating the sampling distribution** for these statistics. What information might we want to know about?

- Variance

Using the distribution

We are **estimating the sampling distribution** for these statistics. What information might we want to know about?

- Variance
- Quantiles

Using the distribution

We are **estimating the sampling distribution** for these statistics. What information might we want to know about?

- Variance
- Quantiles
- Bias (would need to know true parameter) and MSE

Using the distribution

We are **estimating the sampling distribution** for these statistics. What information might we want to know about?

- Variance
- Quantiles
- Bias (would need to know true parameter) and MSE

Using the distribution

We are **estimating the sampling distribution** for these statistics. What information might we want to know about?

- Variance
- Quantiles
- Bias (would need to know true parameter) and MSE

We'll use the first two in order to **develop confidence intervals** and return to bias and MSE later.

Bootstrap Confidence Intervals

Bootstrap confidence intervals

Our goal is to estimate some θ using $\hat{\theta} = T(X_1, X_2, \dots, X_n)$ (we get this without the bootstrap).

Bootstrap confidence intervals

Our goal is to estimate some θ using $\hat{\theta} = T(X_1, X_2, \dots, X_n)$ (we get this without the bootstrap).

We use the bootstrap to understand the **sampling distribution** of $\hat{\theta}$. We can use this to create **confidence intervals** in three basic ways:

- **Large sample approximation** intervals $\hat{\theta} \pm z_{\alpha/2} \hat{\sigma}_{\hat{\theta}}$

Bootstrap confidence intervals

Our goal is to estimate some θ using $\hat{\theta} = T(X_1, X_2, \dots, X_n)$ (we get this without the bootstrap).

We use the bootstrap to understand the **sampling distribution** of $\hat{\theta}$. We can use this to create **confidence intervals** in three basic ways:

- **Large sample approximation** intervals $\hat{\theta} \pm z_{\alpha/2} \hat{\sigma}_{\hat{\theta}}$
- Combining the **estimator on the original sample** with the **bootstrap distribution** in two different ways.

Large sample (Normal approximation)

Many statistics have the property that in **large samples**,

$$T \sim N(\theta, \sigma_T^2) \quad (\text{approx.})$$

Large sample (Normal approximation)

Many statistics have the property that in **large samples**,

$$T \sim N(\theta, \sigma_T^2) \quad (\text{approx.})$$

Using **test inversion** we could create a confidence interval as:

$$T \pm z_{\alpha/2} \sigma_T$$

Large sample (Normal approximation)

Many statistics have the property that in **large samples**,

$$T \sim N(\theta, \sigma_T^2) \quad (\text{approx.})$$

Using **test inversion** we could create a confidence interval as:

$$T \pm z_{\alpha/2} \sigma_T$$

Even if this is known, it may be difficult to compute or estimate σ_T . We substitute the **standard deviation of the bootstrap distribution**:

$$T \pm z_{\alpha} \sqrt{\frac{1}{B-1} \sum_{i=j}^B (T_j^* - \bar{T}^*)^2}$$

Bootstrap intervals vs. usual t -intervals

We derived the following interval:

$$T \pm z_{\alpha} \sqrt{\frac{1}{B-1} \sum_{i=j}^B (T_j^* - \bar{T}^*)^2}$$

Bootstrap intervals vs. usual t -intervals

We derived the following interval:

$$T \pm z_{\alpha} \sqrt{\frac{1}{B-1} \sum_{i=j}^B (T_j^* - \bar{T}^*)^2}$$

Notice that this interval differs from the usual t based interval in three ways:

- We used a Normal quantile instead of a t -quantile. We could have used a t -quantile, but with large B it is basically equivalent.

Bootstrap intervals vs. usual t -intervals

We derived the following interval:

$$T \pm z_{\alpha} \sqrt{\frac{1}{B-1} \sum_{i=j}^B (T_j^* - \bar{T}^*)^2}$$

Notice that this interval differs from the usual t based interval in three ways:

- We used a Normal quantile instead of a t -quantile. We could have used a t -quantile, but with large B it is basically equivalent.
- We used T instead of \bar{T}^* as the center of the interval (observed estimator instead of mean of the bootstrap distribution).

Bootstrap intervals vs. usual t -intervals

We derived the following interval:

$$T \pm z_{\alpha} \sqrt{\frac{1}{B-1} \sum_{i=j}^B (T_j^* - \bar{T}^*)^2}$$

Notice that this interval differs from the usual t based interval in three ways:

- We used a Normal quantile instead of a t -quantile. We could have used a t -quantile, but with large B it is basically equivalent.
- We used T instead of \bar{T}^* as the center of the interval (observed estimator instead of mean of the bootstrap distribution).
- We are estimating the variance of T directly, so there is no $1/\sqrt{n}$ term (variance of sample mean is $\text{Var}(\bar{X}) = (1/n)\text{Var}(X)$).

Normal interval for systolic BP

The trimmed mean is a good example of a statistic that has an asymptotic Normal distribution (see Stigler (1973)) but a tricky variance.

```
> observed_trim + c(1, -1) * qnorm(0.025) * sd(bootstrap_trims)
```

```
[1] 125.6 126.9
```

Basic intervals

As before let $T = T(X_1, X_2, \dots, X_n)$ (the **observed statistic**) and $T^* = T(X_1^*, X_2^*, \dots, X_n^*)$ (the **bootstrap statistic**).

Basic intervals

As before let $T = T(X_1, X_2, \dots, X_n)$ (the **observed statistic**) and $T^* = T(X_1^*, X_2^*, \dots, X_n^*)$ (the **bootstrap statistic**).

Suppose we want to create intervals of the form, $[T - a, T + b]$ so that

$$P(T - a < \theta < T + b) \geq 1 - \alpha$$

(i.e., $[T - a, T + b]$ is a valid CI).

Basic intervals

As before let $T = T(X_1, X_2, \dots, X_n)$ (the **observed statistic**) and $T^* = T(X_1^*, X_2^*, \dots, X_n^*)$ (the **bootstrap statistic**).

Suppose we want to create intervals of the form, $[T - a, T + b]$ so that

$$P(T - a < \theta < T + b) \geq 1 - \alpha$$

(i.e., $[T - a, T + b]$ is a valid CI).

Rewriting, we have the equivalent formulation:

$$P(-b \leq T - \theta \leq a) \geq 1 - \alpha$$

To pick $-b$ and a , we need the distribution of $T - \theta$.

Estimating $T - \theta$ using the bootstrap

The main goal of bootstrapping is **approximate the distribution of T using T^*** . So that

$$P(-b \leq T - \theta \leq a) \approx P(-b \leq T^* - \theta \leq a)$$

Estimating $T - \theta$ using the bootstrap

The main goal of bootstrapping is **approximate the distribution of T using T^*** . So that

$$P(-b \leq T - \theta \leq a) \approx P(-b \leq T^* - \theta \leq a)$$

We know the distribution of T^* (more properly, have an MC estimate of it), **what about θ** ?

Estimating $T - \theta$ using the bootstrap

The main goal of bootstrapping is **approximate the distribution of T using T^*** . So that

$$P(-b \leq T - \theta \leq a) \approx P(-b \leq T^* - \theta \leq a)$$

We know the distribution of T^* (more properly, have an MC estimate of it), **what about θ** ?

Provided T is a good estimator of θ (i.e., $T \xrightarrow{P} \theta$), then

$$P(-b \leq T^* - \theta \leq a) \approx P(-b \leq T^* - T \leq a)$$

Estimating $T - \theta$ using the bootstrap

The main goal of bootstrapping is **approximate the distribution of T using T^*** . So that

$$P(-b \leq T - \theta \leq a) \approx P(-b \leq T^* - \theta \leq a)$$

We know the distribution of T^* (more properly, have an MC estimate of it), **what about θ** ?

Provided T is a good estimator of θ (i.e., $T \xrightarrow{P} \theta$), then

$$P(-b \leq T^* - \theta \leq a) \approx P(-b \leq T^* - T \leq a)$$

This implies that we can pick $-b = T_{\alpha/2}^* - T$ and $a = T_{1-\alpha/2}^* - T$.

Putting it all together

We wanted CIs of the form:

$$P(\theta \in [T - a, T + b]) \geq 1 - \alpha$$

Putting it all together

We wanted CIs of the form:

$$P(\theta \in [T - a, T + b]) \geq 1 - \alpha$$

We found $a = T_{1-\alpha/2}^* - T$ and $b = T - T_{\alpha/2}^*$.

Putting it all together

We wanted CIs of the form:

$$P(\theta \in [T - a, T + b]) \geq 1 - \alpha$$

We found $a = T_{1-\alpha/2}^* - T$ and $b = T - T_{\alpha/2}^*$.

$$[T - a, T + b] = [2T - T_{1-\alpha/2}^*, 2T - T_{\alpha/2}^*]$$

We call these **basic bootstrap confidence intervals**.

Applying to previous examples

```
> ba_trim <- quantile(bootstrap_trims, c(0.975, 0.025))  
> basic_trim <- 2 * observed_trim - ba_trim  
> names(basic_trim) <- NULL # quantile adds some names, we don't need them  
> basic_trim  
  
[1] 125.7 126.9
```


Percentile Intervals

During the previous approach, we were looking for a and b such that

$$P(-b \leq T - \theta \leq a) \geq 1 - \alpha$$

and we had the approximation:

$$P(T_{\alpha/2}^* - T \leq T - \theta \leq T_{1-\alpha/2}^* - T)$$

Percentile Intervals

During the previous approach, we were looking for a and b such that

$$P(-b \leq T - \theta \leq a) \geq 1 - \alpha$$

and we had the approximation:

$$P(T_{\alpha/2}^* - T \leq T - \theta \leq T_{1-\alpha/2}^* - T) = P(T_{\alpha/2}^* \leq 2T - \theta \leq T_{1-\alpha/2}^*)$$

Percentile Intervals

During the previous approach, we were looking for a and b such that

$$P(-b \leq T - \theta \leq a) \geq 1 - \alpha$$

and we had the approximation:

$$P(T_{\alpha/2}^* - T \leq T - \theta \leq T_{1-\alpha/2}^* - T) = P(T_{\alpha/2}^* \leq 2T - \theta \leq T_{1-\alpha/2}^*)$$

Again, since $T \xrightarrow{P} \theta$, $2T \approx 2\theta$, so

$$P(T_{\alpha/2}^* \leq 2T - \theta \leq T_{1-\alpha/2}^*) \approx P(T_{\alpha/2}^* \leq 2\theta - \theta \leq T_{1-\alpha/2}^*) = P(T_{\alpha/2}^* \leq \theta \leq T_{1-\alpha/2}^*)$$

which suggests intervals:

$$[T_{\alpha/2}^*, T_{1-\alpha/2}^*]$$

Percentile Intervals

During the previous approach, we were looking for a and b such that

$$P(-b \leq T - \theta \leq a) \geq 1 - \alpha$$

and we had the approximation:

$$P(T_{\alpha/2}^* - T \leq T - \theta \leq T_{1-\alpha/2}^* - T) = P(T_{\alpha/2}^* \leq 2T - \theta \leq T_{1-\alpha/2}^*)$$

Again, since $T \xrightarrow{P} \theta$, $2T \approx 2\theta$, so

$$P(T_{\alpha/2}^* \leq 2T - \theta \leq T_{1-\alpha/2}^*) \approx P(T_{\alpha/2}^* \leq 2\theta - \theta \leq T_{1-\alpha/2}^*) = P(T_{\alpha/2}^* \leq \theta \leq T_{1-\alpha/2}^*)$$

which suggests intervals:

$$[T_{\alpha/2}^*, T_{1-\alpha/2}^*]$$

These **percentile intervals** also often form the basis for other refinements.

Trimmed mean example

```
> quantile(bootstrap_trims, c(0.025, 0.975))
```

```
 2.5% 97.5%  
125.6 126.8
```

Looking across the three methods

We've see **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.

Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.

Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.
- If you are going to compute the full bootstrap, **why use the Normal approximation?**

Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.
- If you are going to compute the full bootstrap, **why use the Normal approximation?**
- The basic and percentile intervals will always have **the same width**.

Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.
- If you are going to compute the full bootstrap, **why use the Normal approximation?**
- The basic and percentile intervals will always have **the same width**.
- All methods are **approximations** based on the assumptions that:

Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.
- If you are going to compute the full bootstrap, **why use the Normal approximation?**
- The basic and percentile intervals will always have **the same width**.
- All methods are **approximations** based on the assumptions that:
 - $T \approx \theta$ (n is large)

Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.
- If you are going to compute the full bootstrap, **why use the Normal approximation?**
- The basic and percentile intervals will always have **the same width**.
- All methods are **approximations** based on the assumptions that:
 - $T \approx \theta$ (n is large)
 - $\hat{F} \approx F$ (n is large)

Looking across the three methods

We've seen **three methods of creating confidence intervals**. Which one should you use?

- All are equally easy to compute.
- In practice, rarely much difference.
- If you are going to compute the full bootstrap, **why use the Normal approximation?**
- The basic and percentile intervals will always have **the same width**.
- All methods are **approximations** based on the assumptions that:
 - $T \approx \theta$ (n is large)
 - $\hat{F} \approx F$ (n is large)
 - $F^* \approx F$ (B is large)

Using R's `boot` package

Much of the code for running bootstrap estimates is repetitive. Let's use the `boot` package instead.

Using R's `boot` package

Much of the code for running bootstrap estimates is repetitive. Let's use the `boot` package instead.

We need to set up our test statistic function to take a copy of the data and an index for the particular bootstrap sample.

```
> trimmed_mean_boot <- function(x, index, p = 0.1) {  
+   trimmed_mean(x[index], p = p)  
+ }  
> library(boot)  
> boot_tm <- boot(sys_mean, statistic = trimmed_mean_boot, p = 0.2, R = B)
```

```
> boot.ci(boot_tm, type = c("norm", "basic", "perc"))
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 10000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_tm, type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	(125.7, 126.9)	(125.7, 127.0)	(125.5, 126.8)


Calculations and Intervals on Original Scale

Estimating the correlation between systolic and diastolic BP

Research question: Are systolic and diastolic pressure linearly related?

Estimating the correlation between systolic and diastolic BP


Research question: Are systolic and diastolic pressure linearly related?



Bootstrapping the Correlation Coefficient

```
> cor_boot_stat <- function(x, index) {  
+   cor(x[index, 1], x[index, 2], use = "complete")  
+ }  
> library(boot)  
> boot_cor <- boot(nhanes[, c("sys_mean", "dia_mean")],  
+   statistic = cor_boot_stat,  
+   R = 1000)
```

Correlation coefficient distribution



```
> boot.ci(boot_cor, type = c("norm", "basic", "perc"))
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_cor, type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	(0.2269, 0.3117)	(0.2266, 0.3118)	(0.2262, 0.3114)

Calculations and Intervals on Original Scale

Suppose $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ (independent). Then

$$W = \frac{\bar{X} - \mu}{(S^2/n)^{1/2}}$$

has a Student's t -distribution with $n - 1$ degrees of freedom.

Suppose $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ (independent). Then

$$W = \frac{\bar{X} - \mu}{(S^2/n)^{1/2}}$$

has a Student's t -distribution with $n - 1$ degrees of freedom.

More generally we say that a statistic is **studentized** if we subtract off a hypothesized location parameter and divide by an estimate of the standard deviation of the estimator.

Bootstrap-t (percentile) confidence intervals

We noted that **percentile** confidence intervals are frequently the basis of **improved confidence intervals**.

Bootstrap-t (percentile) confidence intervals

We noted that **percentile** confidence intervals are frequently the basis of **improved confidence intervals**.

Define the “studentized” bootstrap replicate

$$W^* = \frac{T^* - T}{\hat{\sigma}^*}$$

Bootstrap-t (percentile) confidence intervals

We noted that **percentile** confidence intervals are frequently the basis of **improved confidence intervals**.

Define the “studentized” bootstrap replicate

$$W^* = \frac{T^* - T}{\hat{\sigma}^*}$$

Undo the studentization to get back to the T scale:

$$[T - \hat{\sigma} W_{1-\alpha/2}^*, T - \hat{\sigma} W_{\alpha/2}^*]$$

Variance Estimators

In the previous algorithm, we used **two different variance estimators** (for notational ease, I'm going to write these using standard deviations instead):

Variance Estimators


In the previous algorithm, we used **two different variance estimators** (for notational ease, I'm going to write these using standard deviations instead):

- $\hat{\sigma}^*$: estimates $\text{Var}(T^*)^{1/2}$ based on a particular bootstrap sample
- $\hat{\sigma}$: estimates $\text{Var}(T)^{1/2}$ based on the original sample

For either of these we could use

- Analytical estimator (e.g., $\text{Var}(\bar{X}) = (1/n)\text{Var}(X)$ and estimate $\text{Var}(X)$ using sample variance statistic S_x^2)
- Bootstrap estimate of variance (“nested bootstrap”)
- The Jackknife (which we'll discuss a bit later)

Log Ratio of Systolic to Diastolic



Bootstrapping the mean

```
> library(boot)
> mean_boot <- function(x, index) { mean(x[index]) }
> boot_mean <- boot(log(sysdia_ratio), statistic = mean_boot, R = 1000)
```

```
> boot.ci(boot_mean, type = c("norm", "basic", "perc"))
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_mean, type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	(0.5931, 0.6086)	(0.5929, 0.6085)	(0.5935, 0.6090)

Calculations and Intervals on Original Scale

Bootstrap-t: sample variance estimator

```
> B <- 1000  
> lsdr <- log(sysdia_ratio)  
> n <- length(lsdr)  
> est_t <- mean(lsdr)  
> est_var_t <- var(lsdr) / n
```



```
> boot_sv <- replicate(B, {  
+   xstar <- sample(lsdrr, replace = TRUE)  
+   (mean(xstar) - est_t) / sqrt(var(xstar) / n)  
+ })  
> (boot_ci_svar <- est_t - sqrt(est_var_t) *  
+   quantile(boot_sv, c(0.975, 0.025)))  
  
      97.5%      2.5%  
0.5932055 0.6088884
```

Bootstrap-t: Nested bootstrap

```
> boot_for_var <- replicate(100, {  
+   xstar <- sample(lsd, replace = TRUE)  
+   mean(xstar)  
+ })  
> boot_var_est <- var(boot_for_var)
```

```

> boot_boot <- replicate(100, {
+   xstar <- sample(lsdrr, replace = TRUE)
+   xstar_boot <- replicate(100, {
+     xstarstar <- sample(xstar, replace = TRUE)
+     mean(xstarstar)
+   })
+   (mean(xstar) - est_t) / sd(xstar_boot)
+ })
> (boot_ci_boot <- est_t - sqrt(boot_var_est) *
+   quantile(boot_boot, c(0.975, 0.025)))

```

```

      97.5%      2.5%
0.5942140 0.6081809

```

Using the boot package

If we return **two values**, the boot package will treat the first as T^* and the second as $\hat{\sigma}_*^2$.

```
> mean_var <- function(x, index) {  
+   xstar <- x[index]  
+   n <- length(xstar)  
+   c(mean(xstar), var(xstar) / n)  
+ }  
> boot_both <- boot(lsd, mean_var, R = 1000)
```

```
> boot.ci(boot_both, type = 'stud')
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_both, type = "stud")
```

Intervals :

Level	Studentized
-------	-------------

95%	(0.5927, 0.6089)
-----	--------------------

Calculations and Intervals on Original Scale

Comparing CIs

	Low	High	Rel. Width
Basic	0.592299625	0.608321349	1.000000000
Percentile	0.593614827	0.609636551	1.000000000
Studentized	0.592678187	0.608857458	1.009833294

Nested bootstrap of the median

With long tailed data (like the log-ratio were using), **the median** may be a better measure of **central tendency** than the mean.

Nested bootstrap of the median

With long tailed data (like the log-ratio were using), **the median** may be a better measure of **central tendency** than the mean.

We can't use sample variance estimate for the median, so we'll use **nested bootstrap**.

```
> median_idx <- function(data, idx) {  
+   median(data[idx])  
+ }  
  
> median_nested <- function(data, idx) {  
+   xstar <- data[idx]  
+   meds <- boot(xstar, median_idx, R = 100)$t # just the T values  
+   return(c(median(xstar), var(meds)))  
+ }
```


Bootstrapping with Parallel Library

```
> library(parallel)
> cl <- makeCluster(detectCores())
> ## load the nested bootstrap components on the cluster
> ignore <- clusterEvalQ(cl, library(boot))
> clusterExport(cl, c("median_idx", "median_nested"))
> boot_median <- boot(lsdrr, median_nested, R = 1000,
+   parallel = "snow", cl = cl, ncpus = detectCores())
> stopCluster(cl)
```

Confidence Intervals

```
> boot.ci(boot_median, type = "stud")
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_median, type = "stud")
```

Intervals :

Level	Studentized
-------	-------------

95%	(0.5423, 0.5550)
-----	--------------------

Calculations and Intervals on Original Scale

Assessing CI Coverage

Confidence Coefficient

For a **parameter** θ and the **random interval** $[A, B]$ we define the **confidence coefficient** c as

$$c = P(A \leq \theta, B \geq \theta)$$

Confidence Coefficient

For a **parameter** θ and the **random interval** $[A, B]$ we define the **confidence coefficient** c as

$$c = P(A \leq \theta, B \geq \theta)$$

If $c \geq 1 - \alpha$ then $[A, B]$ is a valid $(1 - \alpha) \times 100\%$ **Confidence Interval**.

Confidence Coefficient

For a **parameter** θ and the **random interval** $[A, B]$ we define the **confidence coefficient** c as

$$c = P(A \leq \theta, B \geq \theta)$$

If $c \geq 1 - \alpha$ then $[A, B]$ is a valid $(1 - \alpha) \times 100\%$ **Confidence Interval**.

It can be the case that a procedure we claim is a valid CI has $c < 1 - \alpha$. We need to evaluate our procedures to make sure this doesn't happen.

Large Sample Intervals

For

$$X \sim F, E(X) = \mu, \text{Var}(X) < \infty$$

the **central limit theorem** suggests that \bar{X} is approximately

$$\bar{X} \sim N(\mu, \text{Var}(X) / n)$$

Large Sample Intervals

For

$$X \sim F, E(X) = \mu, \text{Var}(X) < \infty$$

the **central limit theorem** suggests that \bar{X} is approximately

$$\bar{X} \sim N(\mu, \text{Var}(X) / n)$$

This suggests that the usual t -intervals are **approximate confidence intervals**:

$$P(\bar{X} - t_{1-\alpha/2}s/\sqrt{n} \leq \mu \leq \bar{X} + t_{1-\alpha/2}s/\sqrt{n}) \approx 1 - \alpha$$

Quality of Approximation

As we saw in HW3, the **quality of this approximation depends on F** , the distribution of X .

Quality of Approximation

As we saw in HW3, the **quality of this approximation depends on F** , the distribution of X .

We investigated the following when $\mu = 1/2$ and found

- $n = 20, X_i \sim \text{Laplace}(1/2)$: good approximation
- $n = 20, X_i \sim \text{Exp}(2)$: poor approximation
- $n = 500, X_i \sim \text{Exp}(2)$: good approximation

Quality of Approximation

As we saw in HW3, the **quality of this approximation depends on F** , the distribution of X .

We investigated the following when $\mu = 1/2$ and found

- $n = 20, X_i \sim \text{Laplace}(1/2)$: good approximation
- $n = 20, X_i \sim \text{Exp}(2)$: poor approximation
- $n = 500, X_i \sim \text{Exp}(2)$: good approximation

Difficult problem: in general, is the Normal approximation valid?

Quick review of HW3 Results: Laplace

```
> k <- 10000
> laplace_intervals <- rerun(k, {
+   t.test(rlaplace(20, mean = 1/2), conf.level = 0.95)$conf.int
+ })
> a1 <- map_dbl(laplace_intervals, ~ .x[1] <= 1/2 & .x[2] >= 1/2)
> binom.test(sum(a1), k, conf.level = 0.99)$conf.int

[1] 0.9439078 0.9552592
attr(,"conf.level")
[1] 0.99
```

Quick review of HW3 Results: Exponential

```
> exponential_intervals <- rerun(k, {  
+   t.test(rexp(20, rate = 2), conf.level = 0.95)$conf.int  
+ })  
> a2 <- map_dbl(exponential_intervals, ~ .x[1] <= 1/2 & .x[2] >= 1/2)  
> binom.test(sum(a2), k, conf.level = 0.99)$conf.int  
  
[1] 0.9132797 0.9273164  
attr(,"conf.level")  
[1] 0.99
```

Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

- Generate sample n from known distribution

Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

- Generate sample n from known distribution
- Perform bootstrap procedure B times

Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

- Generate sample n from known distribution
- Perform bootstrap procedure B times
- Compute Normal theory, basic and percentile intervals

Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

- Generate sample n from known distribution
- Perform bootstrap procedure B times
- Compute Normal theory, basic and percentile intervals
- Repeat k times

Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

- Generate sample n from known distribution
- Perform bootstrap procedure B times
- Compute Normal theory, basic and percentile intervals
- Repeat k times

Evaluating Bootstrap Procedure

Goal: Do the bootstrap confidence interval procedures have proper confidence coefficients?

- Generate sample n from known distribution
- Perform bootstrap procedure B times
- Compute Normal theory, basic and percentile intervals
- Repeat k times

Note, requires about $n \times B \times k$ steps. Let's **parallelize!**

Setting up local cluster

```
> library(boot)
> library(parallel)
> cl <- makeCluster(detectCores() - 1)
> ## Example usage:
> mean_boot <- function(x, idx) { mean(x[idx]) }
> a <- boot(rlaplace(20, 1/2), mean_boot, R = 1000,
+         parallel = "snow",
+         cl = cl,
+         ncpus = 2) %>% # bug: must be greater than 1
+         boot.ci(type = c("norm", "basic", "perc"))
```

Formatted results

```
> print(a)
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = ., type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	(0.3568, 1.1857)	(0.3439, 1.1986)	(0.3674, 1.2221)

Calculations and Intervals on Original Scale

What is the structure of `boot.ci` result?

```
> names(a)

[1] "R"          "t0"          "call"        "normal"      "basic"       "percent"

> a$normal ; a$basic ; a$percent

      conf
[1,] 0.95 0.3567675 1.185686

      conf
[1,] 0.95 975.98 25.03 0.3439484 1.198577

      conf
[1,] 0.95 25.03 975.98 0.367446 1.222074
```

Pulling out just the intervals

```
> getCIs <- function(boot_ci_result) {  
+   with(boot_ci_result,  
+       matrix(c(normal[2:3],  
+               basic[4:5],  
+               percent[4:5]), nrow = 2))  
+ }  
> getCIs(a)
```

	[,1]	[,2]	[,3]
[1,]	0.3567675	0.3439484	0.367446
[2,]	1.1856857	1.1985765	1.222074

Putting it together: Laplace

```
> k <- 1000 ; R <- 1000
> laplace_bootstrap_cis <- rerun(k, {
+   boot(rlaplace(20, 1/2), mean_boot, R = R,
+       parallel = "snow", cl = cl, ncpus = 2) %>%
+   boot.ci(type = c("norm", "basic", "perc")) %>%
+   getCIs
+ })
```

Putting it together: Exponential

```
> exp_bootstrap_cis <- rerun(k, {  
+   boot(rexp(20, 2), mean_boot, R = R,  
+       parallel = "snow", cl = cl, ncpus = 2) %>%  
+       boot.ci(type = c("norm", "basic", "perc")) %>%  
+       getCIs  
+ })
```

Counting Covering CIs

```
> covers <- function(x) { x[1, ] <= 1/2 & x[2, ] >= 1/2 }  
> laplace_bootstrap_covers <- map(laplace_bootstrap_cis, covers) %>%  
+   simplify2array %>% rowSums  
> exp_bootstrap_covers <- map(exp_bootstrap_cis, covers) %>%  
+   simplify2array %>% rowSums
```

Coverage Rates: Laplace

```
> coverage_ci <- function(x) { binom.test(x, k, conf.level = 0.99)$conf.int  
> map(laplace_bootstrap_covers, coverage_ci)
```

```
[[1]]
```

```
[1] 0.9099948 0.9517603
```

```
[[2]]
```

```
[1] 0.9237124 0.9619124
```

```
[[3]]
```

```
[1] 0.8975841 0.9422873
```

Coverage Rates: Exponential

```
> map(exp_bootstrap_covers, coverage_ci)
```

```
[[1]]
```

```
[1] 0.8621267 0.9140733
```

```
[[2]]
```

```
[1] 0.8511989 0.9051019
```

```
[[3]]
```

```
[1] 0.8698147 0.9203143
```

Why poor coverage rates?

Remember, we need $\hat{F}(x)$ **to be close to** $F(x)$ for all x .

Why poor coverage rates?

Remember, we need $\hat{F}(x)$ **to be close to** $F(x)$ for all x .

We had a sample size of 20. Is this enough for a good approximation? (Apparently not.)


Why poor coverage rates?

Remember, we need $\hat{F}(x)$ **to be close to** $F(x)$ for all x .

We had a sample size of 20. Is this enough for a good approximation? (Apparently not.)

Let's repeat this procedure at **increasingly large sample sizes** (and compare to `t.test` along the way).

Plotting results



Interpreting Coverage Rates

In the previous examples, we saw that **all** of the methods had **confidence coefficients** below the targeted level.

Interpreting Coverage Rates

In the previous examples, we saw that **all** of the methods had **confidence coefficients** below the targeted level.

For the t-test based intervals, we know this is because **the CLT Normal approximation only holds as $n \rightarrow \infty$.**

Interpreting Coverage Rates

In the previous examples, we saw that **all** of the methods had **confidence coefficients** below the targeted level.

For the t-test based intervals, we know this is because **the CLT Normal approximation only holds as $n \rightarrow \infty$** .

For the other methods, we only have that $\hat{F}(x) \xrightarrow{P} F$, another **law of large numbers** type result. The bootstrap requires the approximation $\hat{F}(x) \approx F$ to hold, which for small samples is tenuous.

Interpreting Coverage Rates

In the previous examples, we saw that **all** of the methods had **confidence coefficients** below the targeted level.

For the t-test based intervals, we know this is because **the CLT Normal approximation only holds as $n \rightarrow \infty$** .

For the other methods, we only have that $\hat{F}(x) \xrightarrow{P} F$, another **law of large numbers** type result. The bootstrap requires the approximation $\hat{F}(x) \approx F$ to hold, which for small samples is tenuous.

Remember: the primary benefit of the bootstrap is **trading Monte Carlo methods for analytic methods** not magically making new data.

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.
- Main idea: If $X_i \stackrel{\text{iid}}{\sim} F$, **estimate** \hat{F} in order to apply **inversion method**:
 $\hat{Q}(U) \sim F$ (approximately).

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.
- Main idea: If $X_i \stackrel{\text{iid}}{\sim} F$, **estimate \hat{F}** in order to apply **inversion method**:
 $\hat{Q}(U) \sim F$ (approximately).
- The **bootstrap sample** of $X_1^*, X_2^*, \dots, X_n^*$ is used to compute
 $T^* = T(X_1^*, \dots, X_n^*)$.

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.
- Main idea: If $X_i \stackrel{\text{iid}}{\sim} F$, **estimate** \hat{F} in order to apply **inversion method**:
 $\hat{Q}(U) \sim F$ (approximately).
- The **bootstrap sample** of $X_1^*, X_2^*, \dots, X_n^*$ is used to compute
 $T^* = T(X_1^*, \dots, X_n^*)$.
- Since $T^* \sim T$ (approx.) we can create confidence intervals:

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.
- Main idea: If $X_i \stackrel{\text{iid}}{\sim} F$, **estimate** \hat{F} in order to apply **inversion method**:
 $\hat{Q}(U) \sim F$ (approximately).
- The **bootstrap sample** of $X_1^*, X_2^*, \dots, X_n^*$ is used to compute
 $T^* = T(X_1^*, \dots, X_n^*)$.
- Since $T^* \sim T$ (approx.) we can create confidence intervals:
 - **Large sample**: $T \pm z_{\alpha/2} \sigma_{T^*}$

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.
- Main idea: If $X_i \stackrel{\text{iid}}{\sim} F$, **estimate** \hat{F} in order to apply **inversion method**: $\hat{Q}(U) \sim F$ (approximately).
- The **bootstrap sample** of $X_1^*, X_2^*, \dots, X_n^*$ is used to compute $T^* = T(X_1^*, \dots, X_n^*)$.
- Since $T^* \sim T$ (approx.) we can create confidence intervals:
 - **Large sample**: $T \pm z_{\alpha/2} \sigma_{T^*}$
 - **Basic**: $[2T - T_{1-\alpha/2}^*, 2T - T_{\alpha/2}^*]$

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.
- Main idea: If $X_i \stackrel{\text{iid}}{\sim} F$, **estimate \hat{F}** in order to apply **inversion method**:
 $\hat{Q}(U) \sim F$ (approximately).
- The **bootstrap sample** of $X_1^*, X_2^*, \dots, X_n^*$ is used to compute
 $T^* = T(X_1^*, \dots, X_n^*)$.
- Since $T^* \sim T$ (approx.) we can create confidence intervals:
 - **Large sample**: $T \pm z_{\alpha/2} \sigma_{T^*}$
 - **Basic**: $[2T - T_{1-\alpha/2}^*, 2T - T_{\alpha/2}^*]$
 - **Percentile**: $[T_{\alpha/2}^*, T_{1-\alpha/2}^*]$

Summary

- Goal: **estimate** θ using an **estimator** $\hat{\theta} = T = T(X_1, X_2, \dots, X_n)$.
- Need to know the **sampling distribution of $\hat{\theta}$** to find **confidence intervals**.
- Main idea: If $X_i \stackrel{\text{iid}}{\sim} F$, **estimate** \hat{F} in order to apply **inversion method**:
 $\hat{Q}(U) \sim F$ (approximately).
- The **bootstrap sample** of $X_1^*, X_2^*, \dots, X_n^*$ is used to compute
 $T^* = T(X_1^*, \dots, X_n^*)$.
- Since $T^* \sim T$ (approx.) we can create confidence intervals:
 - **Large sample**: $T \pm z_{\alpha/2} \sigma_{T^*}$
 - **Basic**: $[2T - T_{1-\alpha/2}^*, 2T - T_{\alpha/2}^*]$
 - **Percentile**: $[T_{\alpha/2}^*, T_{1-\alpha/2}^*]$
 - **Studentized**: $[T - \sigma W_{1-\alpha/2}^*, T - \sigma W_{\alpha/2}^*]$, where $W^* = (T^* - T)/\hat{\sigma}^*$ (requires within bootstrap sample variance estimate)